

```
In [1]: import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.model_selection import train_test_split
```

```
In [2]: df=pd.read_csv("/Users/bob/Downloads/FP1_air/csvs_per_year/csvs_per_
df
```

Out[2]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	
0	2011-11-01 01:00:00	NaN	1.0	NaN	NaN	154.0	84.0	NaN	NaN	NaN	6.0	NaN	I
1	2011-11-01 01:00:00	2.5	0.4	3.5	0.26	68.0	92.0	3.0	40.0	24.0	9.0	1.54	
2	2011-11-01 01:00:00	2.9	NaN	3.8	NaN	96.0	99.0	NaN	NaN	NaN	NaN	NaN	
3	2011-11-01 01:00:00	NaN	0.6	NaN	NaN	60.0	83.0	2.0	NaN	NaN	NaN	NaN	I
4	2011-11-01 01:00:00	NaN	NaN	NaN	NaN	44.0	62.0	3.0	NaN	NaN	3.0	NaN	I
...	
209923	2011-09-01 00:00:00	NaN	0.2	NaN	NaN	5.0	19.0	44.0	NaN	NaN	NaN	NaN	I
209924	2011-09-01 00:00:00	NaN	0.1	NaN	NaN	6.0	29.0	NaN	11.0	NaN	7.0	NaN	I
209925	2011-09-01 00:00:00	NaN	NaN	NaN	0.23	1.0	21.0	28.0	NaN	NaN	NaN	1.44	I
209926	2011-09-01 00:00:00	NaN	NaN	NaN	NaN	3.0	15.0	48.0	NaN	NaN	NaN	NaN	I
209927	2011-09-01 00:00:00	NaN	NaN	NaN	NaN	4.0	33.0	38.0	13.0	NaN	NaN	NaN	I

209928 rows × 14 columns

In [3]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 209928 entries, 0 to 209927
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
---  --
 0   date        209928 non-null object  
 1   BEN         51393 non-null  float64
 2   CO          87127 non-null  float64
 3   EBE         51350 non-null  float64
 4   NMHC        43517 non-null  float64
 5   NO          208954 non-null float64
 6   NO_2        208973 non-null float64
 7   O_3         122049 non-null float64
 8   PM10        103743 non-null float64
 9   PM25        51079 non-null  float64
10   SO_2        87131 non-null  float64
11   TCH         43519 non-null  float64
12   TOL         51175 non-null  float64
13   station     209928 non-null int64  
dtypes: float64(12), int64(1), object(1)
memory usage: 22.4+ MB
```

```
In [4]: df1=df.dropna()
df1
```

Out [4]:


	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TO
1	2011-11-01 01:00:00	2.5	0.4	3.5	0.26	68.0	92.0	3.0	40.0	24.0	9.0	1.54	8.
6	2011-11-01 01:00:00	0.7	0.3	1.1	0.16	17.0	66.0	7.0	22.0	16.0	2.0	1.36	1.
25	2011-11-01 02:00:00	1.8	0.3	2.8	0.20	34.0	76.0	3.0	34.0	21.0	8.0	1.71	7.
30	2011-11-01 02:00:00	1.0	0.4	1.3	0.18	31.0	67.0	5.0	25.0	18.0	3.0	1.40	2.
49	2011-11-01 03:00:00	1.3	0.2	2.4	0.22	29.0	72.0	3.0	33.0	20.0	8.0	1.75	6.
...
209862	2011-08-31 22:00:00	0.4	0.1	1.0	0.06	1.0	13.0	33.0	21.0	6.0	5.0	1.26	0.
209881	2011-08-31 23:00:00	0.9	0.1	1.8	0.16	11.0	45.0	30.0	32.0	17.0	3.0	1.34	4.
209886	2011-08-31 23:00:00	0.6	0.1	1.1	0.05	1.0	12.0	48.0	19.0	7.0	5.0	1.26	0.
209905	2011-09-01 00:00:00	0.6	0.1	1.3	0.15	6.0	35.0	34.0	21.0	12.0	3.0	1.32	3.
209910	2011-09-01 00:00:00	0.7	0.1	1.1	0.04	1.0	12.0	46.0	8.0	5.0	5.0	1.25	0.

16460 rows × 14 columns

```
In [5]: df1=df1.drop(["date"],axis=1)
```


```
In [6]: sns.heatmap(df1.corr())
```

```
Out[6]: <Axes: >
```



```
In [7]: plt.plot(df1["EBE"],df1["NMHC"],"o")
```

```
Out[7]: [matplotlib.lines.Line2D at 0x7fcc2d3e080>]
```



```
In [8]: data=df[["EBE","NMHC"]]
```

```
In [9]: x=df1.drop(["EBE"],axis=1)
y=df1["EBE"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```


Linear

```
In [10]: li=LinearRegression()
li.fit(x_train,y_train)
```

```
Out[10]: ▼ LinearRegression
LinearRegression()
```

```
In [11]: prediction=li.predict(x_test)
plt.scatter(y_test,prediction)
```

```
Out[11]: <matplotlib.collections.PathCollection at 0x7fccb1523fd0>
```



```
In [12]: lis=li.score(x_test,y_test)
```

```
In [13]: df1["TCH"].value_counts()
```

```
Out[13]: 1.30      897
         1.29      878
         1.28      856
         1.31      827
         1.27      820
         ...
         2.89         1
         3.06         1
         3.36         1
         2.99         1
         3.49         1
         Name: TCH, Length: 171, dtype: int64
```

```
In [14]: df1.loc[df1["TCH"]<1.40,"TCH"]=1  
df1.loc[df1["TCH"]>1.40,"TCH"]=2  
df1["TCH"].value_counts()
```

```
Out[14]: 1.0    12828  
        2.0     3632  
        Name: TCH, dtype: int64
```

Lasso

```
In [15]: la=Lasso(alpha=5)  
la.fit(x_train,y_train)
```

```
Out[15]: 

▼



Lasso




Lasso(alpha=5)


```

```
In [16]: prediction1=la.predict(x_test)  
plt.scatter(y_test,prediction1)
```

```
Out[16]: <matplotlib.collections.PathCollection at 0x7fccb159c430>
```



```
In [17]: las=la.score(x_test,y_test)
```

Ridge

```
In [18]: rr=Ridge(alpha=1)
         rr.fit(x_train,y_train)
```


```
Out[18]:
```

▼ Ridge

Ridge(alpha=1)

```
In [19]: prediction2=rr.predict(x_test)
         plt.scatter(y_test,prediction2)
```

```
Out[19]: <matplotlib.collections.PathCollection at 0x7fcca30222c0>
```



```
In [20]: rrs=rr.score(x_test,y_test)
```


ElasticNet


```
In [21]: en=ElasticNet()  
en.fit(x_train,y_train)
```

```
Out[21]: ▼ ElasticNet  
ElasticNet()
```

```
In [22]: prediction2=rr.predict(x_test)  
plt.scatter(y_test,prediction2)
```

```
Out[22]: <matplotlib.collections.PathCollection at 0x7fcca2fdb850>
```



```
In [23]: ens=en.score(x_test,y_test)
```

```
In [24]: print(rr.score(x_test,y_test))  
rr.score(x_train,y_train)
```

```
0.8279312677210368
```

```
Out[24]: 0.8141050323866895
```

Logistic

```
In [25]: g={"TCH":{1.0:"Low",2.0:"High"}}
df1=df1.replace(g)
df1["TCH"].value_counts()
```

```
Out[25]: Low      12828
         High      3632
         Name: TCH, dtype: int64
```

```
In [26]: x=df1.drop(["TCH"],axis=1)
y=df1["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [27]: lo=LogisticRegression()
lo.fit(x_train,y_train)
```

```
Out[27]: ▾ LogisticRegression
         LogisticRegression()
```

```
In [28]: prediction3=lo.predict(x_test)
plt.scatter(y_test,prediction3)
```

```
Out[28]: <matplotlib.collections.PathCollection at 0x7fcca2251600>
```



```
In [29]: los=lo.score(x_test,y_test)
```

Random Forest

```
In [30]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
```

```
In [31]: g1={"TCH":{"Low":1.0,"High":2.0}}
df1=df1.replace(g1)
```

```
In [32]: x=df1.drop(["TCH"],axis=1)
y=df1["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [33]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out [33]: ▼ RandomForestClassifier
RandomForestClassifier()
```

```
In [34]: parameter={
    'max_depth': [1,2,4,5,6],
    'min_samples_leaf': [5,10,15,20,25],
    'n_estimators': [10,20,30,40,50]
}
```

```
In [35]: grid_search=GridSearchCV(estimator=rfc,param_grid=parameter,cv=2,sc
grid_search.fit(x_train,y_train)
```

```
Out [35]: ▶ GridSearchCV
▶ estimator: RandomForestClassifier
    ▶ RandomForestClassifier
```

```
In [36]: rfcs=grid_search.best_score_
```

```
In [37]: rfc_best=grid_search.best_estimator_
```

```
In [38]: from sklearn.tree import plot_tree

plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_name=
Text(0.5350877192982456, 0.35714285714285715, '0_3 <= 5.5\ngini =
0.489\nsamples = 276\nvalue = [246, 182]\nclass = Yes'),
Text(0.5175438596491229, 0.21428571428571427, 'NMHC <= 0.205\ngini
i = 0.427\nsamples = 92\nvalue = [47, 105]\nclass = No'),
Text(0.5087719298245614, 0.07142857142857142, 'gini = 0.499\nsampl
les = 41\nvalue = [33, 30]\nclass = Yes'),
Text(0.5263157894736842, 0.07142857142857142, 'gini = 0.265\nsampl
les = 51\nvalue = [14, 75]\nclass = No'),
Text(0.5526315789473685, 0.21428571428571427, 'PM10 <= 17.5\ngini
= 0.402\nsamples = 184\nvalue = [199, 77]\nclass = Yes'),
Text(0.543859649122807, 0.07142857142857142, 'gini = 0.064\nsampl
es = 19\nvalue = [29, 1]\nclass = Yes'),
Text(0.5614035087719298, 0.07142857142857142, 'gini = 0.427\nsampl
les = 165\nvalue = [170, 76]\nclass = Yes'),
Text(0.6052631578947368, 0.35714285714285715, 'NO <= 82.5\ngini =
0.439\nsamples = 167\nvalue = [89, 184]\nclass = No'),
Text(0.5877192982456141, 0.21428571428571427, 'NO_2 <= 67.5\ngini
= 0.426\nsamples = 153\nvalue = [76, 171]\nclass = No'),
Text(0.5789473684210527, 0.07142857142857142, 'gini = 0.113\nsampl
les = 49\nvalue = [5, 78]\nclass = No'),
```

```
In [39]: print("Linear:",lis)
print("Lasso:",las)
print("Ridge:",rrs)
print("ElasticNet:",ens)
print("Logistic:",los)
print("Random Forest:",rfcs)
```

```
Linear: 0.8279633125878545
Lasso: 0.5823338875996914
Ridge: 0.8279312677210368
ElasticNet: 0.7138887521604254
Logistic: 0.7806804374240583
Random Forest: 0.8900364520048603
```

Best Model is Random Forest

```
In [40]: df2=pd.read_csv("/Users/bob/Downloads/FP1_air/csvs_per_year/csvs_pe
df2
```

Out[40]:

	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TI
0	2012-09-01 01:00:00	NaN	0.2	NaN	NaN	7.0	18.0	NaN	NaN	NaN	2.0	NaN	N
1	2012-09-01 01:00:00	0.3	0.3	0.7	NaN	3.0	18.0	55.0	10.0	9.0	1.0	NaN	z
2	2012-09-01 01:00:00	0.4	NaN	0.7	NaN	2.0	10.0	NaN	NaN	NaN	NaN	NaN	.
3	2012-09-01 01:00:00	NaN	0.2	NaN	NaN	1.0	6.0	50.0	NaN	NaN	NaN	NaN	N
4	2012-09-01 01:00:00	NaN	NaN	NaN	NaN	1.0	13.0	54.0	NaN	NaN	3.0	NaN	N
...
210715	2012-03-01 00:00:00	NaN	0.6	NaN	NaN	37.0	84.0	14.0	NaN	NaN	NaN	NaN	N
210716	2012-03-01 00:00:00	NaN	0.4	NaN	NaN	5.0	76.0	NaN	17.0	NaN	7.0	NaN	N
210717	2012-03-01 00:00:00	NaN	NaN	NaN	0.34	3.0	41.0	24.0	NaN	NaN	NaN	1.34	N
210718	2012-03-01 00:00:00	NaN	NaN	NaN	NaN	2.0	44.0	36.0	NaN	NaN	NaN	NaN	N
210719	2012-03-01 00:00:00	NaN	NaN	NaN	NaN	2.0	56.0	40.0	18.0	NaN	NaN	NaN	N

210720 rows × 14 columns

In [41]: df2.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 210720 entries, 0 to 210719
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   date        210720 non-null  object
1   BEN         51511 non-null   float64
2   CO          87097 non-null   float64
3   EBE         51482 non-null   float64
4   NMHC        30736 non-null   float64
5   NO          209871 non-null  float64
6   NO_2        209872 non-null  float64
7   O_3         122339 non-null  float64
8   PM10        104838 non-null  float64
9   PM25        52164 non-null   float64
10  SO_2        87333 non-null   float64
11  TCH         30736 non-null   float64
12  TOL         51373 non-null   float64
13  station     210720 non-null  int64
dtypes: float64(12), int64(1), object(1)
memory usage: 22.5+ MB
```

```
In [42]: df3=df2.dropna()  
df3
```

Out[42]:


	date	BEN	CO	EBE	NMHC	NO	NO_2	O_3	PM10	PM25	SO_2	TCH	TO
6	2012-09-01 01:00:00	0.4	0.2	0.8	0.24	1.0	7.0	57.0	11.0	7.0	2.0	1.33	0.
30	2012-09-01 02:00:00	0.4	0.2	0.7	0.24	1.0	5.0	55.0	5.0	5.0	2.0	1.33	0.
54	2012-09-01 03:00:00	0.4	0.2	0.7	0.24	1.0	4.0	56.0	6.0	4.0	2.0	1.33	0.
78	2012-09-01 04:00:00	0.3	0.2	0.7	0.25	1.0	5.0	54.0	6.0	5.0	2.0	1.34	0.
102	2012-09-01 05:00:00	0.4	0.2	0.7	0.24	1.0	3.0	53.0	8.0	5.0	2.0	1.33	0.
...
210654	2012-02-29 22:00:00	0.6	0.3	0.5	0.09	1.0	35.0	57.0	25.0	21.0	3.0	1.12	2.
210673	2012-02-29 23:00:00	2.0	0.4	2.4	0.21	16.0	79.0	20.0	37.0	25.0	12.0	1.33	6.
210678	2012-02-29 23:00:00	0.7	0.3	0.6	0.09	1.0	27.0	63.0	22.0	18.0	3.0	1.11	1.
210697	2012-03-01 00:00:00	1.5	0.4	1.7	0.21	16.0	79.0	17.0	28.0	21.0	11.0	1.34	4.
210702	2012-03-01 00:00:00	0.6	0.3	0.5	0.09	1.0	23.0	61.0	18.0	16.0	3.0	1.11	1.

10916 rows × 14 columns

```
In [43]: df3=df3.drop(["date"],axis=1)
```

```
In [44]: sns.heatmap(df3.corr())
```

```
Out[44]: <Axes: >
```



```
In [45]: x=df3.drop(["TCH"],axis=1)
y=df3["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear


```
In [46]: li=LinearRegression()
li.fit(x_train,y_train)
```

```
Out[46]: ▼ LinearRegression
LinearRegression()
```



```
In [47]: prediction=li.predict(x_test)
plt.scatter(y_test,prediction)
```

```
Out[47]: <matplotlib.collections.PathCollection at 0x7fccc36d3010>
```



```
In [48]: lis=li.score(x_test,y_test)
```

```
In [49]: df3["TCH"].value_counts()
```

```
Out[49]: 1.30    737
         1.31    676
         1.32    644
         1.33    552
         1.29    529
         ...
         2.39     1
         2.20     1
         2.72     1
         3.11     1
         2.70     1
         Name: TCH, Length: 167, dtype: int64
```

```
In [50]: df3.loc[df3["TCH"]<1.40,"TCH"]=1  
df3.loc[df3["TCH"]>1.40,"TCH"]=2  
df3["TCH"].value_counts()
```

```
Out[50]: 1.0    8772  
        2.0    2144  
        Name: TCH, dtype: int64
```

```
In [ ]:
```

Lasso

```
In [51]: la=Lasso(alpha=5)  
la.fit(x_train,y_train)
```

```
Out[51]: 

▼



Lasso




Lasso(alpha=5)


```

```
In [52]: prediction1=la.predict(x_test)  
plt.scatter(y_test,prediction1)
```

```
Out[52]: <matplotlib.collections.PathCollection at 0x7fccc3747370>
```



```
In [53]: las=la.score(x_test,y_test)
```

Ridge


```
In [54]: rr=Ridge(alpha=1)  
rr.fit(x_train,y_train)
```

```
Out [54]:
```

▼	Ridge
	Ridge(alpha=1)

```
In [55]: prediction2=rr.predict(x_test)  
plt.scatter(y_test,prediction2)
```

```
Out [55]: <matplotlib.collections.PathCollection at 0x7fccc37477f0>
```



```
In [56]: rrs=rr.score(x_test,y_test)
```


ElasticNet

```
In [57]: en=ElasticNet()  
en.fit(x_train,y_train)
```

```
Out[57]: ▾ ElasticNet  
ElasticNet()
```

```
In [58]: prediction2=rr.predict(x_test)  
plt.scatter(y_test,prediction2)
```

```
Out[58]: <matplotlib.collections.PathCollection at 0x7fccb2023b50>
```



```
In [59]: ens=en.score(x_test,y_test)
```

```
In [60]: print(rr.score(x_test,y_test))  
rr.score(x_train,y_train)
```

```
0.7027769553239515
```

```
Out[60]: 0.6820735566707625
```

Logistic

```
In [61]: g={"TCH":{1.0:"Low",2.0:"High"}}
df3=df3.replace(g)
df3["TCH"].value_counts()
```

```
Out[61]: Low      8772
         High     2144
         Name: TCH, dtype: int64
```

```
In [62]: x=df3.drop(["TCH"],axis=1)
y=df3["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [63]: lo=LogisticRegression()
lo.fit(x_train,y_train)
```

```
Out[63]: ▼ LogisticRegression
LogisticRegression()
```

```
In [64]: prediction3=lo.predict(x_test)
plt.scatter(y_test,prediction3)
```

```
Out[64]: <matplotlib.collections.PathCollection at 0x7fccc36a6020>
```



```
In [65]: los=lo.score(x_test,y_test)
```

Random Forest

```
In [66]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
```

```
In [67]: g1={"TCH":{"Low":1.0,"High":2.0}}
df3=df3.replace(g1)
```

```
In [68]: x=df3.drop(["TCH"],axis=1)
y=df3["TCH"]
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

```
In [69]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out [69]: ▼ RandomForestClassifier
RandomForestClassifier()
```

```
In [70]: parameter={
    'max_depth': [1,2,4,5,6],
    'min_samples_leaf': [5,10,15,20,25],
    'n_estimators': [10,20,30,40,50]
}
```

```
In [71]: grid_search=GridSearchCV(estimator=rfc,param_grid=parameter,cv=2,sc
grid_search.fit(x_train,y_train)
```

```
Out [71]: ► GridSearchCV
► estimator: RandomForestClassifier
    ► RandomForestClassifier
```

```
In [72]: rfcs=grid_search.best_score_
```

```
In [73]: rfc_best=grid_search.best_estimator_
```

```
In [74]: from sklearn.tree import plot_tree

plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_nam
```

```
Out [74]: [Text(0.5377604166666666, 0.9285714285714286, 'NMHC <= 0.275\ngini
= 0.315\nsamples = 4831\nvalue = [6147, 1494]\nclass = Yes'),
Text(0.28385416666666667, 0.7857142857142857, 'CO <= 0.35\ngini =
0.132\nsamples = 4074\nvalue = [6002, 460]\nclass = Yes'),
Text(0.15104166666666666, 0.6428571428571429, 'PM25 <= 11.5\ngini
= 0.084\nsamples = 3548\nvalue = [5413, 248]\nclass = Yes'),
Text(0.08333333333333333, 0.5, 'NO_2 <= 13.5\ngini = 0.046\nsampl
es = 2726\nvalue = [4256, 103]\nclass = Yes'),
Text(0.041666666666666664, 0.35714285714285715, 'PM10 <= 10.5\ngi
ni = 0.01\nsamples = 1437\nvalue = [2290, 12]\nclass = Yes'),
Text(0.020833333333333332, 0.21428571428571427, 'NMHC <= 0.265\ng
ini = 0.003\nsamples = 806\nvalue = [1291, 2]\nclass = Yes'),
Text(0.010416666666666666, 0.07142857142857142, 'gini = 0.0\nsamp
les = 787\nvalue = [1265, 0]\nclass = Yes'),
Text(0.03125, 0.07142857142857142, 'gini = 0.133\nsamples = 19\nv
alue = [26, 2]\nclass = Yes'),
Text(0.0625, 0.21428571428571427, 'EBE <= 2.2\ngini = 0.02\nsampl
es = 631\nvalue = [999, 10]\nclass = Yes'),
Text(0.052083333333333336, 0.07142857142857142, 'gini = 0.012\nsa
mples = 626\nvalue = [804, 6]\nclass = Yes')]
```

```
In [75]: print("Linear:",lis)
print("Lasso:",las)
print("Ridge:",rrs)
print("ElasticNet:",ens)
print("Logistic:",los)
print("Random Forest:",rfcs)
```

```
Linear: 0.7025995094680152
Lasso: -7.83493655487355e-06
Ridge: 0.7027769553239515
ElasticNet: 0.38196670444638303
Logistic: 0.8003053435114503
Random Forest: 0.9339087106113775
```

Best model is Random Forest

